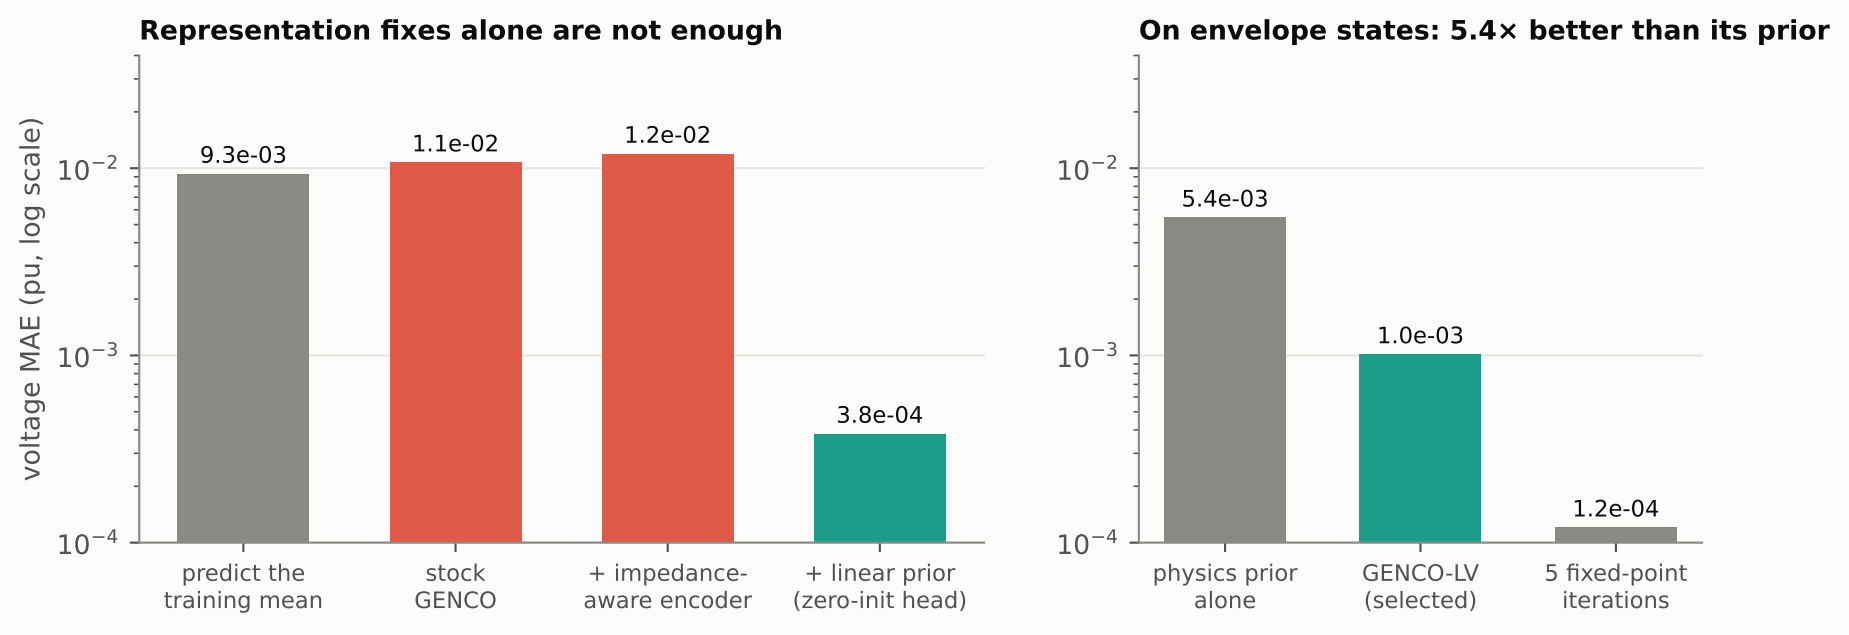
Representation fixes alone are not enough
voltage MAE (pu, log scale)
10^{-2}
10^{-3}
10^{-4}
9.3e-03
1.1e-02
1.2e-02
3.8e-04
predict the training mean
stock GENCO
+ impedance-aware encoder
+ linear prior (zero-init head)
On envelope states: 5.4× better than its prior
5.4e-03
1.0e-03
1.2e-04
physics prior alone
GENCO-LV (selected)
5 fixed-point iterations